

```

import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University, and
modified for ece239as at UCLA.
"""

class KNN(object):

    def __init__(self):
        pass

    def train(self, X, y):
        """
        Inputs:
        - X is a numpy array of size (num_examples, D)--data
        - y is a numpy array of size (num_examples, )--label
        """
        self.X_train = X
        self.y_train = y

    def compute_distances(self, X, norm=None):
        """
        Compute the distance between each test point in X and each training point
        in self.X_train.

        Inputs:
        - X: A numpy array of shape (num_test, D) containing test data.
        - norm: the function with which the norm is taken.

        Returns:
        - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
            is the Euclidean distance between the ith test point and the jth training
            point.
        """
        if norm is None:
            norm = lambda x: np.sqrt(np.sum(x**2))
            #norm = 2

        num_test = X.shape[0]
        num_train = self.X_train.shape[0]
        dists = np.zeros((num_test, num_train))

        for i in np.arange(num_test):
            for j in np.arange(num_train):
                # ===== #
                # YOUR CODE HERE:
                #   Compute the distance between the ith test point and the jth
                #   training point using norm(), and store the result in dists[i, j].
                # ===== #
                dists[i,j] = norm(X[i]-self.X_train[j])
                pass
                # ===== #
                # END YOUR CODE HERE
                # ===== #

```

```

    return dists

def compute_L2_distances_vectorized(self, X):
    """
    Compute the distance between each test point in X and each training point
    in self.X_train WITHOUT using any for loops.

    Inputs:
    - X: A numpy array of shape (num_test, D) containing test data.

    Returns:
    - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
      is the Euclidean distance between the ith test point and the jth training
      point.
    """
    num_test = X.shape[0]
    num_train = self.X_train.shape[0]
    dists = np.zeros((num_test, num_train))

    # ===== #
    # YOUR CODE HERE:
    # Compute the L2 distance between the ith test point and the jth
    # training point and store the result in dists[i, j]. You may
    # NOT use a for loop (or list comprehension). You may only use
    # numpy operations.
    #
    # HINT: use broadcasting. If you have a shape (N,1) array and
    # a shape (M,) array, adding them together produces a shape (N, M)
    # array.
    # ===== #

    Xtr = np.sum(self.X_train**2, axis = 1)          # (5000, 3072) -->
    (5000, ^2) --> (5000, sum) --> (5000,)
    Xte = np.sum(X**2, axis = 1).reshape(X.shape[0], 1) # (500, 3072) -->
    (500, ^2) --> (500, sum) --> (500,) --> (500, 1)
    dists = np.sqrt(Xtr + Xte - 2*X.dot(self.X_train.T)) # (5000,)+(500,1) ---
    >(500, 5000)
    #pass

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return dists

def predict_labels(self, dists, k=1):
    """
    Given a matrix of distances between test points and training points,
    predict a label for each test point.

    Inputs:
    - dists: A numpy array of shape (num_test, num_train) where dists[i, j]
      gives the distance between the ith test point and the jth training point.

```

Returns:

- y: A numpy array of shape (num_test,) containing predicted labels for the test data, where y[i] is the predicted label for the test point X[i].

"""

```
num_test = dists.shape[0]
y_pred = np.zeros(num_test)
for i in np.arange(num_test):
    # A list of length k storing the labels of the k nearest neighbors to
    # the ith test point.
    closest_y = []
    # ===== #
    # YOUR CODE HERE:
    # Use the distances to calculate and then store the labels of
    # the k-nearest neighbors to the ith test point. The function
    # numpy.argsort may be useful.
    #
    # After doing this, find the most common label of the k-nearest
    # neighbors. Store the predicted label of the ith training example
    # as y_pred[i]. Break ties by choosing the smaller label.
    # ===== #
    neighbors_index = np.argsort(dists[i,:], axis = 0)
    #print (neighbors_index)
    closest_y = self.y_train[neighbors_index[:k]]
    #print(closest_y) #(500,k) with labels
    freq = np.bincount(closest_y)
    y_pred[i] = np.argmax(freq)

    # ===== #
    # END YOUR CODE HERE
    # ===== #
#print(len(y_pred))
return y_pred
```