

```

import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University, and
modified for ece239as at UCLA.
"""
class SVM(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the SVM. Note that it has shape (C, D)
        where C is the number of classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims)

    def loss(self, X, y):
        """
        Calculates the SVM loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 ≤ c < C.

        Returns a tuple of:
        - loss as single float
        """

        # compute the loss and the gradient
        num_classes = self.W.shape[0]
        num_train = X.shape[0]
        loss = 0.0
        #print(self.W.shape)

        # ===== #
        # YOUR CODE HERE:
        # Calculate the normalized SVM loss, and store it as 'loss'.
        # (That is, calculate the sum of the losses of all the training
        # set margins, and then normalize the loss by the number of
        # training examples.)
        # ===== #
        for i in np.arange(num_train):
            for j in np.arange(num_classes):
                if j == y[i]:
                    continue
                aj = self.W[j].dot(X[i])
                ay = self.W[y[i]].dot(X[i])
                loss += max(0, 1 + aj - ay)
        loss = loss / num_train

```

```

# ===== #
# END YOUR CODE HERE
# ===== #

return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
             the gradient of the loss with respect to W.
    """
    # compute the loss and the gradient
    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    loss = 0.0
    delta = 1
    grad = np.zeros_like(self.W)

    # ===== #
    # YOUR CODE HERE:
    #   Calculate the SVM loss and the gradient. Store the gradient in
    #   the variable grad.ears
    # ===== #
    for i in np.arange(num_train):
        for j in np.arange(num_classes):
            if j == y[i]:
                continue
            aj = self.W[j].dot(X[i]) #incorrect class score
            ay = self.W[y[i]].dot(X[i]) #correct class score
            z = delta+aj-ay
            loss += max(0, z)
            if z > 0:
                grad[j,:] +=X[i]
                grad[y[i],:] -=X[i]

    # ===== #
    # END YOUR CODE HERE
    # ===== #
    #print (grad.shape) #(10,3073)
    loss /= num_train
    grad /= num_train

    return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]

```

```

self.W[ix] = oldval + h # increment by h
fxph = self.loss(X, y)
self.W[ix] = oldval - h # decrement by h
fxmh = self.loss(X,y) # evaluate f(x - h)
self.W[ix] = oldval # reset

grad_numerical = (fxph - fxmh) / (2 * h)
grad_analytic = your_grad[ix]
rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) +
    abs(grad_analytic))
print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical,
    grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    delta = 1
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the SVM loss WITHOUT any for loops.
    # ===== #
    scores = np.matmul(self.W,(X.T))
    #print('score shape', scores.shape) #(10,500)
    correct_scores = scores[y,np.arange(num_train)]
    #print('correct score shape', correct_scores.shape) #(500,)
    z = delta+scores-correct_scores
    margins = np.maximum(0, z) #broadcasting
    #print(margins.shape) (10,500)
    margins[y, np.arange(num_train)] = 0 #correct class score rows should be 0
    loss = np.sum(margins)
    loss = loss / num_train
    # ===== #
    # END YOUR CODE HERE
    # ===== #
    # ===== #
    # YOUR CODE HERE:
    # Calculate the SVM grad WITHOUT any for loops.
    # ===== #
    grad_scores = np.zeros(margins.shape)
    grad_scores[margins > 0] = 1 #indicator function gives 1 if z>0, gives 0 if
    <=0
    #print(grad_scores.shape) #(10,500)
    grad_scores[y, np.arange(num_train)] = -np.sum(margins > 0, axis = 0)
    #print (X.shape) (500,3073)
    grad = np.matmul(grad_scores,(X))
    grad = grad / num_train
    # ===== #
    # END YOUR CODE HERE
    # ===== #

```

```

return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
        means that X[i] has label 0 ≤ c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is
        number of classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the
        weights of self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:
        # Sample batch_size elements from the training data for use in
        # gradient descent. After sampling,
        # - X_batch should have shape: (batch_size, dim)
        # - y_batch should have shape: (batch_size,)
        # The indices should be randomly generated to reduce correlations
        # in the dataset. Use np.random.choice. It's okay to sample with
        # replacement.
        # ===== #
        indices = np.random.choice(np.arange(num_train), batch_size)
        #print('indices shape', indices.shape) #(200,)
        X_batch = X[indices,:]
        #print('X_batch shape', X_batch.shape) #(200,3073)
        y_batch = y[indices]
        #print('y_batch shape', y_batch.shape) #(200,)

        # ===== #
        # END YOUR CODE HERE
        # ===== #

```

```

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

# ===== #
# YOUR CODE HERE:
#   Update the parameters, self.W, with a gradient step
# ===== #
self.W -= learning_rate * grad
# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])

    # ===== #
    # YOUR CODE HERE:
    #   Predict the labels given the training data with the parameter self.W.
    # ===== #
    scores = X.dot(self.W.T)
    #print(scores.shape)
    y_pred = np.argmax(scores, axis = 1)
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return y_pred

```